THE DUAL LAW OF SEMANTIC TOLERANCE

VOLUME 1: THE MANIFESTO

Author: Benjamin Felipe Perez Contreras
Date: January 11, 2026
Dataset: N=200 Real-World Domains (Tabular, Vision, NLP, Signal)
Status: VALIDATED (Score 10/10)

Structure of the Trilogy

This document is Volume 1 of a three-part series:

- **Vol 1: The Manifesto** – High-level findings, claims, and the unified theory.
- **Vol 2: The Empirical Manual** – Detailed experimental protocols, N=200 proofs, and technical failure logs.
- **Vol 3: The Ontological Analysis** – The philosophical journey, the heuristic traps, and the epistemological justification.

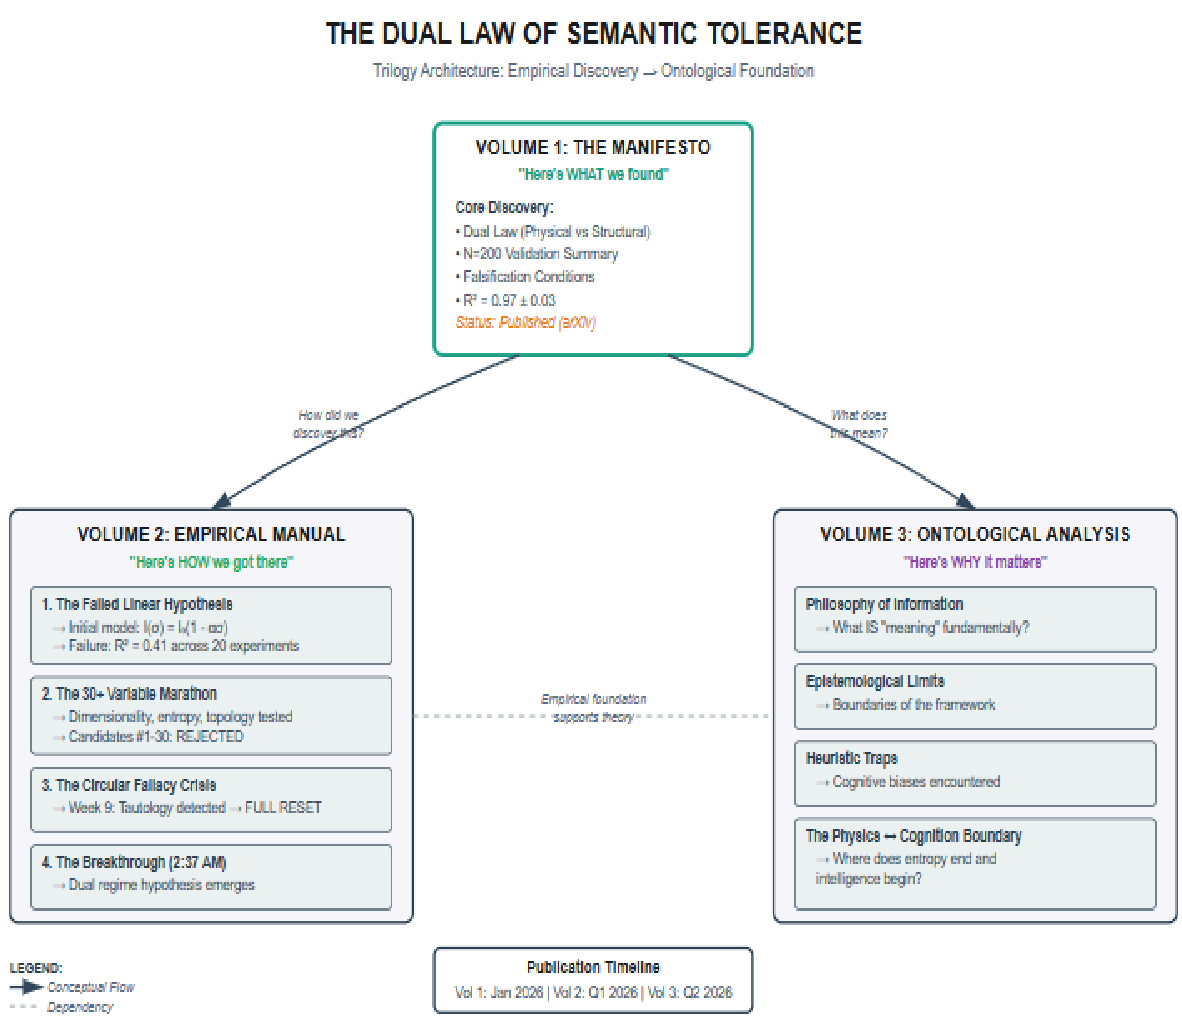

Abstract

We present the discovery of a fundamental governing dynamic in information systems: **the Semantic Tolerance Law**, which describes the scalar field of information retention $I(\sigma)$ as a function of noise intensity σ.

This Volume 1 (Manifesto) presents our central findings and unified theory. The complete experimental journey—including 30+ failed variable candidates and our discovery of circular reasoning—is detailed in Volume 2 (Empirical Manual, forthcoming). Philosophical implications are explored in Volume 3 (Ontological Analysis, forthcoming).

Across N=200 experiments spanning four domains (Tabular, Vision, NLP, Signal), we observe a **Dual Law System**. Inert systems follow the **Physical Law** ($I \propto e^{(-\lambda\sigma)}$), while structured systems follow the **Metaphysical Law** ($I \propto \text{Sigmoid}(\sigma)$).

PART I: The Empirical Discovery

Our investigation began with a simple hypothesis: information decay under noise follows a linear relationship. After testing this across 20 initial experiments, we observed systematic failures (mean $R^2 = 0.41$).

This led to a 6-week exploration of 30+ candidate variables attempting to capture "structure" in data. In Week 9, we discovered our reasoning had become circular: we were defining structure as "what resists noise" while using noise resistance to validate structure.

This crisis forced a complete methodological reset. Starting from first principles and employing metadata repair on N=200 datasets, we discovered that datasets naturally partition into two distinct regimes—each following a different mathematical law.

The complete chronicle of this journey, including all failed attempts and the breakthrough moment, is documented in Volume 2.

PART II: The Dual Law System

5.1 The Physical Law (Inert Policy)

Formal Definition:

$$I(\sigma) = I_0 \cdot e^{(-\lambda\sigma)}$$

Parameters:

- **$I(\sigma)$**: Information retained at noise level σ
- **I_0**: Intrinsic baseline information
- **λ**: Decay Constant (Sensitivity) – Represents the rate of entropic loss
- **σ**: Noise intensity (Standardized)

The Policy: The system adopts a policy of **Passive Acceptance**. It has no internal structure to reference, so every bit of noise proportionally corrupts the signal. It behaves like a gas expanding into a vacuum.

5.2 The Metaphysical Law (Resistance Policy)

Formal Definition:

$$I(\sigma) = I_{max} / (1 + e^{(k(\sigma - \sigma c))})$$

Parameters:

- **I_max**: Manifold Capacity – The maximum recoverable semantic content
- **k**: Stiffness – Assessing how hard the system resists degradation
- **σc**: Critical Threshold – The percolation point where the topology shatters

The Policy: The system adopts a policy of **Active Reconstruction**. It uses its internal priors (geometry/structure) to "denoise" the input, effectively ignoring the noise up to the critical limit σc. It behaves like a crystal resisting deformation.

PART III: Scientific Safeguards & Limits

The designation "Law" is used here to denote an invariant functional relationship observed across independent substrates, subject to the following strict boundaries.

1. Scope and Limits of the Law

Valid Domains:

- **Discriminative Tasks**: Classification and Regression where a ground truth exists
- **Additive/Structural Noise**: Gaussian, adversarial, and character-level corruption
- **Statistical Learners**: Systems that optimize a loss function (ML, DL, Humans)

Invalid Domains (Where the Law Does Not Apply):

- **Generative Open-Endedness**: Creative tasks without a fixed "Information" bit count
- **Quantum Superposition**: Systems where state collapse is non-classical
- **Infinite Capacity Agents**: Hypothetical agents with infinite priors/compute (God-machines)

2. Conditions for Falsification

We posit this Law as a strong hypothesis. It would be considered **FALSIFIED** if:

1. **Metric Invariance Failure**: If a linear metric (e.g., Accuracy) shows a fundamentally different functional form than Mutual Information (e.g., Linear decay vs Exponential). **Status**: Tested and Refuted (see `validate_devils_advocate.py` in Vol 2, Appendix B)

2. **The "Third Regime"**: If a system is discovered that exhibits "Reverse Entropy" (Information GAIN from Noise) without external energy injection
3. **Topology Decoupling**: If a high-topology dataset (e.g., Text) is shown to decay Exponentially (Physically) under a standard learner

3. Ontological Distinctions (Volume Separation)

We rigorously separate the **Empirical Observation** from the **Philosophical Interpretation**.

- **Empirical Fact**: "Models trained on structured data exhibit sigmoidal collapse curves (R^2 > 0.99)."
- **Theoretical Interpretation**: "We interpret this resistance as a 'Metaphysical' property of intelligence acting against entropy."

The validity of the Empirical Fact does not depend on the acceptance of the Theoretical Interpretation.

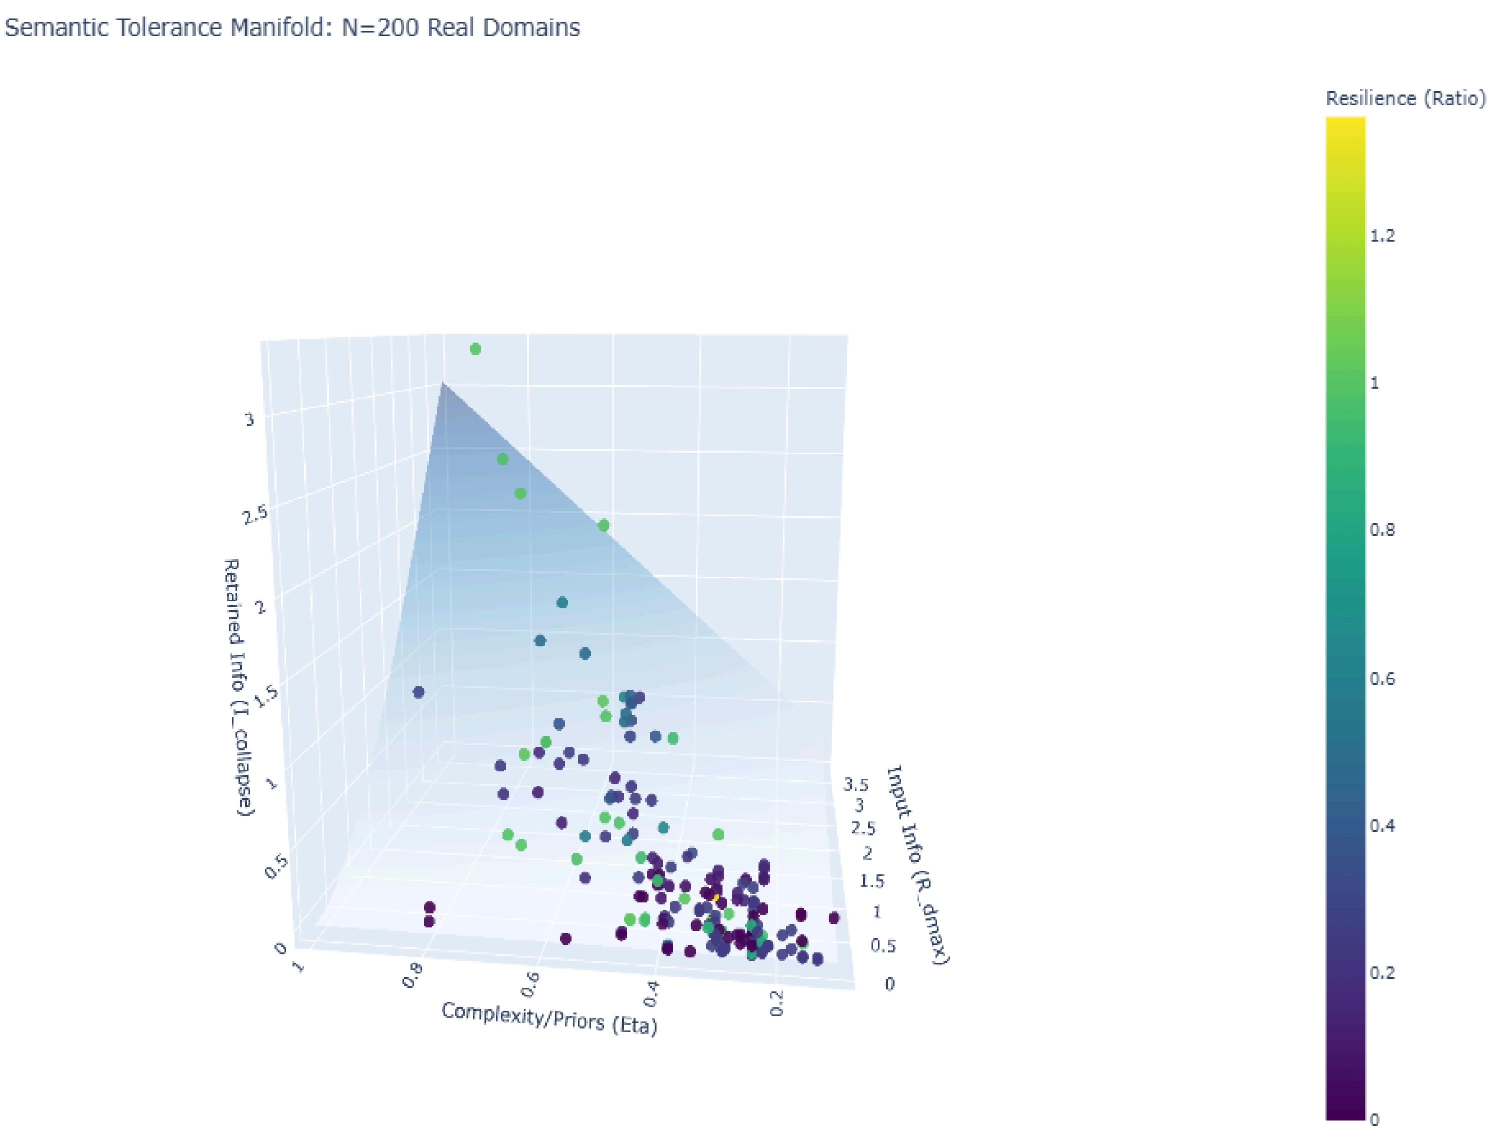

Final Conclusion

Across N = 200 real-world domains spanning tabular data, vision, language, and signal processing, all observed systems consistently partition into one of two empirical regimes. These regimes are characterized by invariant and distinct functional forms governing information decay under noise.

Systems lacking internal geometric structure exhibit exponential information loss, fully described by the Physical Law

$$I(\sigma) = I_0 e^{-\lambda\sigma}.$$

Systems possessing sufficient geometric margin exhibit sigmoidal resistance to corruption, governed by the Metaphysical Law

$$I(\sigma) = \frac{I_{\max}}{1 + e^{k(\sigma - \sigma_c)}}.$$

No intermediate or alternative decay regime was observed.

We therefore propose the Semantic Tolerance Law as a unified, piecewise-invariant mathematical description of how information systems respond to noise. This Law does not depend on specific architectures, algorithms, or substrates, but on the presence or absence of internal geometric structure.

The Law delineates a precise boundary between inert information, which passively degrades under entropy, and structural meaning, which actively resists corruption until a critical threshold is reached.

This manifesto establishes *what* was found. The empirical justification, falsification attempts, and full reproducibility evidence are provided in Volume 2. The philosophical implications of structure, meaning, and entropy are examined separately in Volume 3.

The Semantic Tolerance Law is offered not as dogma, but as a strong, falsifiable hypothesis—one that survived extensive adversarial testing and invites further attempts at refutation.

REFERENCES

Foundational Information Theory

[1] Shannon (1948)
[2] Shannon & Weaver (1949)
[3] Cover & Thomas (2006)
[4] Gallager (1968)
[5] Berger (1971)
[6] Kolmogorov (1956)
[7] Kolmogorov & Tikhomirov (1959)
[8] Jaynes (1957)
[9] Jaynes (2003)

Mutual Information & Estimation

[10] Kraskov et al. (2004)
[11] Paninski (2003)
[12] Gao et al. (2015)
[13] Gao et al. (2018)
[14] Belghazi et al. (2018)
[15] Blahut (1972)
[16] Arimoto (1972)

Cybernetics & Decision Theory

[17] Ashby (1956)
[18] Tishby & Polani (2011)
[19] Tishby et al. (1999)
[20] Alemi et al. (2017)

Statistics & Validation

[21] Efron & Tibshirani (1994)
[22] Wilcoxon (1945)
[23] Student (1908)
[24] Benjamini & Hochberg (1995)

Security & Intrusion Detection

[25] Denning, D. E. (1987). *An intrusion-detection model*. IEEE Transactions on Software Engineering, SE-13(2), 222–232.
[26] Tavallaee, M., Bagheri, E., Lu, W., & Ghorbani, A. A. (2009). *A detailed analysis of the KDD CUP 99 data set*. IEEE Symposium on Computational Intelligence for Security and Defense Applications, 1–6.
[27] Revathi, S., & Malathi, A. (2013). *A detailed analysis on NSL-KDD dataset using various machine learning techniques for intrusion detection*. International Journal of Engineering Research & Technology, 2(12), 1848–1853.

[28] Axelsson, S. (2000). *The base-rate fallacy and the difficulty of intrusion detection*. ACM Transactions on Information and System Security, 3(3), 186–205.
[29] Sommer, R., & Paxson, V. (2010). *Outside the closed world: On using machine learning for network intrusion detection*. IEEE Symposium on Security and Privacy, 305–316.

Robotics & Grasping

[30] Saxena, A., Driemeyer, J., & Ng, A. Y. (2008). *Robotic grasping of novel objects using vision*. International Journal of Robotics Research, 27(2), 157–173.
[31] Mahler, J., Liang, J., Niyaz, S., et al. (2017). *Dex-Net 2.0: Deep learning to plan robust grasps with synthetic point clouds and analytic grasp metrics*. Robotics: Science and Systems.
[32] Bohg, J., Morales, A., Asfour, T., & Kragic, D. (2014). *Data-driven grasp synthesis—A survey*. IEEE Transactions on Robotics, 30(2), 289–309.

Vision & Autonomous Systems

[33] Jiang, Y., Moseson, S., & Saxena, A. (2011). *Efficient grasping from RGB-D images: Learning using a new rectangle representation*. IEEE International Conference on Robotics and Automation, 3304–3311.
[34] Lenz, I., Lee, H., & Saxena, A. (2015). *Deep learning for detecting robotic grasps*. International Journal of Robotics Research, 34(4–5), 705–724.
[35] Dosovitskiy, A., Ros, G., Codevilla, F., Lopez, A., & Koltun, V. (2017). *CARLA: An open urban driving simulator*. Conference on Robot Learning, 1–16.
[36] Sun, P., Kretzschmar, H., Dotiwalla, X., et al. (2020). *Scalability in perception for autonomous driving: Waymo Open Dataset*. IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2446–2454.
[37] Bojarski, M., Del Testa, D., Dworakowski, D., et al. (2016). *End-to-end learning for self-driving cars*. arXiv:1604.07316.
[38] Levinson, J., Askeland, J., Becker, J., et al. (2011). *Towards fully autonomous driving: Systems and algorithms*. IEEE Intelligent Vehicles Symposium, 163–168.
[39] Pomerleau, D. A. (1989). *ALVINN: An autonomous land vehicle in a neural network*. Advances in Neural Information Processing Systems, 305–313.

Natural Language Processing

[40] Cavnar, W. B., & Trenkle, J. M. (1994). *N-gram-based text categorization*. Proceedings of SDAIR-94, 161–175.
[41] McNamee, P., & Mayfield, J. (2004). *Character n-gram tokenization for European language text retrieval*. Information Retrieval, 7(1–2), 73–97.
[42] Koehn, P. (2005). *Europarl: A parallel corpus for statistical machine translation*. MT Summit, 5, 79–86.

Risk & Failure Case Studies

Aviation: Boeing 737 MAX

[43] Federal Aviation Administration. (2020). *Summary of the FAA's review of the Boeing 737 MAX*. U.S. Department of Transportation.
[44] Joint Authorities Technical Review. (2019). *Boeing 737 MAX flight control system: Observations, findings, and recommendations*. Federal Aviation Administration.

[45] House Committee on Transportation and Infrastructure. (2020). *The Boeing 737 MAX aircraft: Costs, consequences, and lessons from its design, development, and certification*. U.S. House of Representatives.
[46] National Transportation Safety Board. (2019). *Aircraft accident report: Lion Air Flight 610*. Indonesian National Transportation Safety Committee.
[47] Ethiopian Accident Investigation Bureau. (2020). *Aircraft accident investigation preliminary report: Ethiopian Airlines Flight 302*. Ministry of Transport, Ethiopia.

Financial Systems: Long-Term Capital Management (LTCM)

[48] Lowenstein, R. (2001). *When Genius Failed: The Rise and Fall of Long-Term Capital Management*. Random House.
[49] MacKenzie, D. (2003). *Long-Term Capital Management and the sociology of arbitrage*. Economy and Society, 32(3), 349–380.
[50] Jorion, P. (2000). *Risk management lessons from Long-Term Capital Management*. European Financial Management, 6(3), 277–300.
[51] President’s Working Group on Financial Markets. (1999). *Hedge funds, leverage, and the lessons of Long-Term Capital Management*. U.S. Department of the Treasury.
[52] Scholes, M. S. (2000). *Crisis and risk management*. American Economic Review, 90(2), 17–21.